

Intro-End2End

March 22, 2020

0.1 Introduction

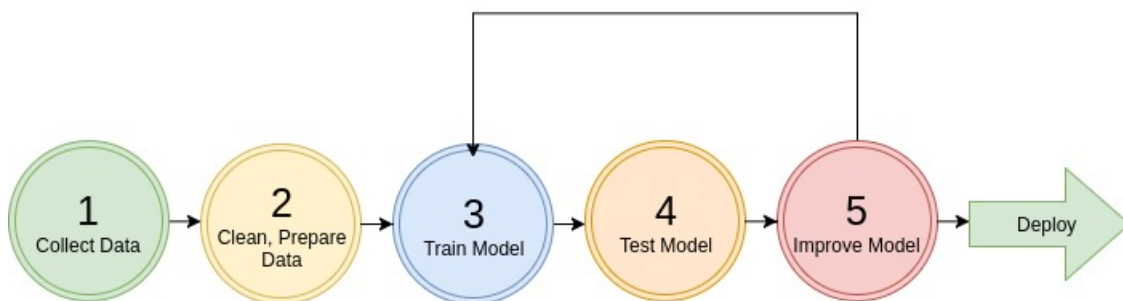
Welcome to **CS188 - Data Science Fundamentals!** We plan on having you go through some grueling training so you can start crunching data out there... in today's day and age "data is the new oil" or perhaps "snake oil" nonetheless, there's a lot of it, each with different purity (so pure that perhaps you could feed off it for a life time) or dirty which then at that point you can either decide to dump it or try to weed out something useful (that's where they need you...)

In this project you will work through an example project end to end.

Here are the main steps:

1. Get the data
2. Visualize the data for insights
3. Preprocess the data for your machine learning algorithm
4. Select a model and train
5. Does it meet the requirements? Fine tune the model

Steps to Machine Learning



0.2 Working with Real Data

It is best to experiment with real-data as opposed to artificial datasets.

There are many different open datasets depending on the type of problems you might be interested in!

Here are a few data repositories you could check out: - [UCI Datasets](#) - [Kaggle Datasets](#) - [AWS Datasets](#)

Below we will run through an California Housing example collected from the 1990's.

0.3 Setup

```
[1]: import sys
assert sys.version_info >= (3, 5) # python>=3.5
import sklearn
assert sklearn.__version__ >= "0.20" # sklearn >= 0.20

import numpy as np #numerical package in python
import os
%matplotlib inline
import matplotlib.pyplot as plt #plotting package

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt

# Where to save the figures
ROOT_DIR = "."
IMAGES_PATH = os.path.join(ROOT_DIR, "images")
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_name, tight_layout=True, fig_extension="png", resolution=300):
    """
        plt.savefig wrapper. refer to
        https://matplotlib.org/3.1.1/api/_as_gen/matplotlib.pyplot.savefig.html
    """
    path = os.path.join(IMAGES_PATH, fig_name + "." + fig_extension)
    print("Saving figure", fig_name)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)

[2]: import os
import tarfile
import urllib
DATASET_PATH = os.path.join("datasets", "housing")
```

0.4 Intro to Data Exploration Using Pandas

In this section we will load the dataset, and visualize different features using different types of plots.

Packages we will use: - **Pandas**: is a fast, flexible and expressive data structure widely used for tabular and multidimensional datasets. - **Matplotlib**: is a 2d python plotting library which you can use to create quality figures (you can plot almost anything if you're willing to code it out!) -

other plotting libraries: `seaborn`, `ggplot2`

```
[3]: import pandas as pd

def load_housing_data(housing_path):
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)

[4]: housing = load_housing_data(DATASET_PATH) # we load the pandas dataframe
housing.head() # show the first few elements of the dataframe
               # typically this is the first thing you do
               # to see how the dataframe looks like
```

```
[4]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

A dataset may have different types of features - real valued - Discrete (integers) - categorical (strings)

The two categorical features are essentially the same as you can always map a categorical string/character to an integer.

In the dataset example, all our features are real valued floats, except ocean proximity which is categorical.

```
[5]: # to see a concise summary of data types, null values, and counts
      # use the info() method on the dataframe
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
```

```

4   total_bedrooms      20433 non-null float64
5   population          20640 non-null float64
6   households          20640 non-null float64
7   median_income       20640 non-null float64
8   median_house_value  20640 non-null float64
9   ocean_proximity     20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

```

[6]: # you can access individual columns similarly
# to accessing elements in a python dict
housing["ocean_proximity"].head() # added head() to avoid printing many columns.
↪ .

```

```

[6]: 0    NEAR BAY
1    NEAR BAY
2    NEAR BAY
3    NEAR BAY
4    NEAR BAY
Name: ocean_proximity, dtype: object

```

```

[7]: # to access a particular row we can use iloc
housing.iloc[1]

```

```

[7]: longitude      -122.22
latitude          37.86
housing_median_age    21
total_rooms          7099
total_bedrooms       1106
population           2401
households           1138
median_income        8.3014
median_house_value   358500
ocean_proximity      NEAR BAY
Name: 1, dtype: object

```

```

[8]: # one other function that might be useful is
# value_counts(), which counts the number of occurrences
# for categorical features
housing["ocean_proximity"].value_counts()

```

```

[8]: <1H OCEAN      9136
INLAND          6551
NEAR OCEAN       2658
NEAR BAY         2290
ISLAND            5
Name: ocean_proximity, dtype: int64

```

```
[9]: # The describe function compiles your typical statistics for each
# column
housing.describe()
```

```
[9]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

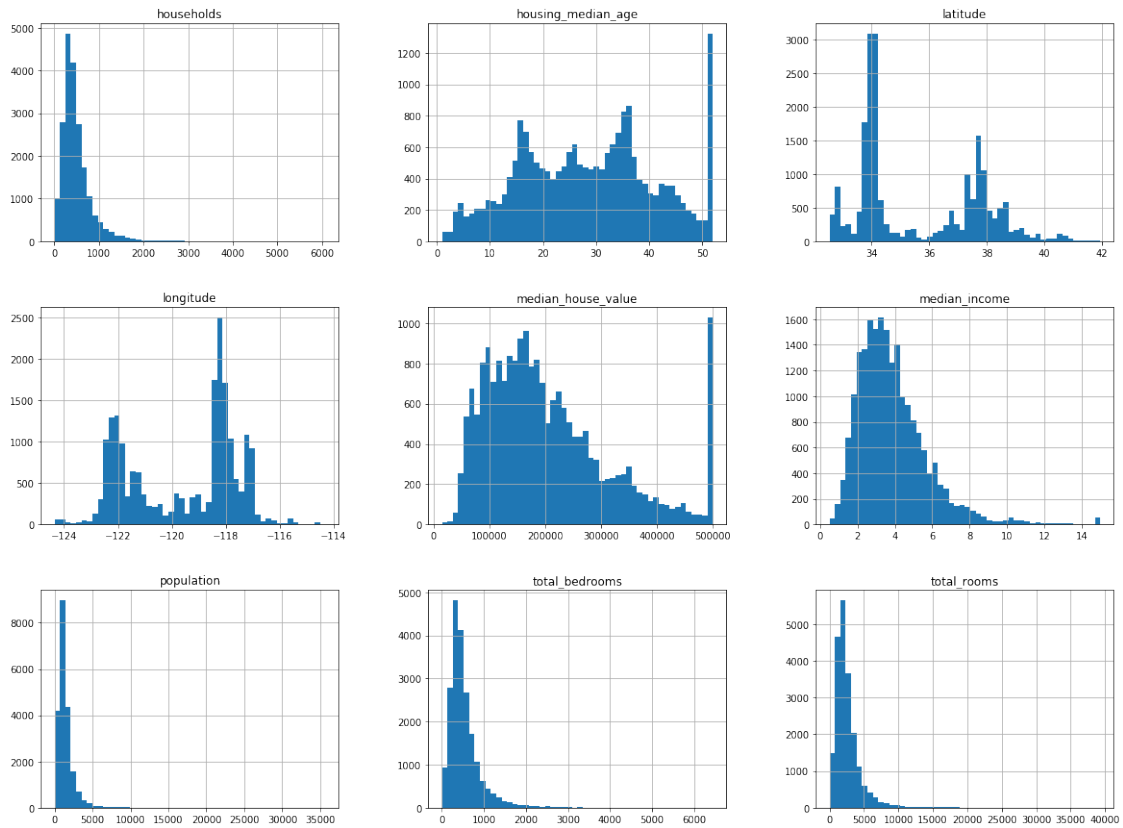
	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

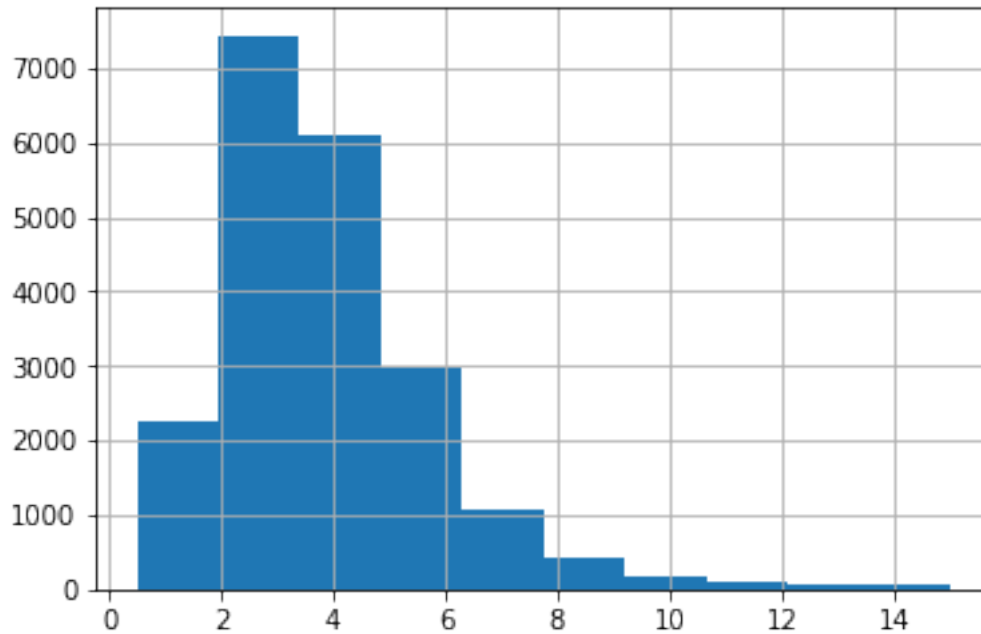
If you want to learn about different ways of accessing elements or other functions it's useful to check out the getting started section [here](#)

0.5 Let's start visualizing the dataset

```
[10]: # We can draw a histogram for each of the dataframes features
# using the hist function
housing.hist(bins=50, figsize=(20,15))
# save_fig("attribute_histogram_plots")
plt.show() # pandas internally uses matplotlib, and to display all the figures
# the show() function must be called
```



```
[11]: # if you want to have a histogram on an individual feature:
housing["median_income"].hist()
plt.show()
```



We can convert a floating point feature to a categorical feature by binning or by defining a set of intervals.

For example, to bin the households based on median_income we can use the `pd.cut` function

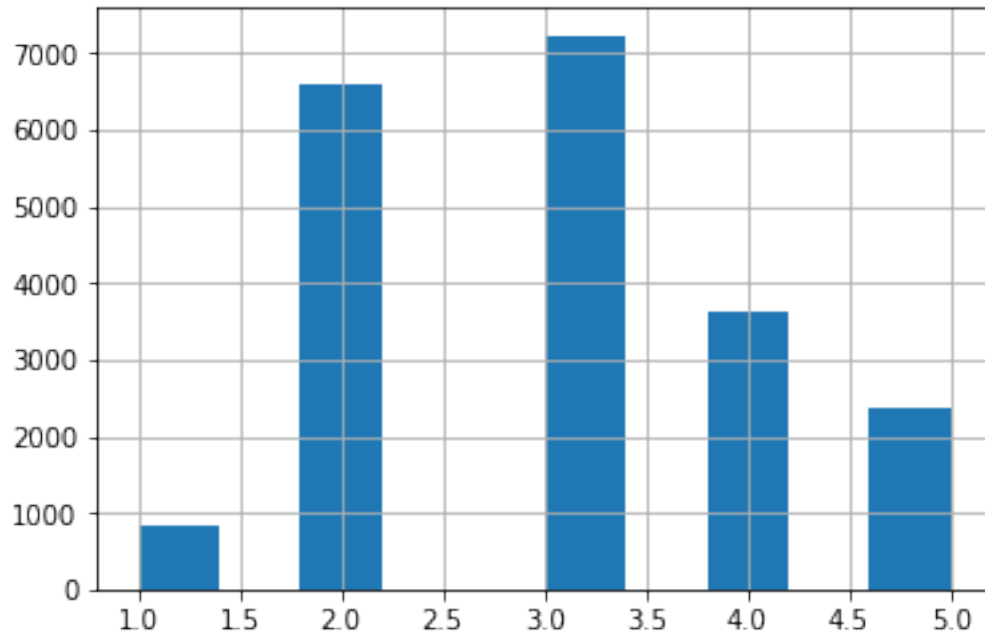
```
[12]: # assign each bin a categorical value [1, 2, 3, 4, 5] in this case.
housing["income_cat"] = pd.cut(housing["median_income"],
                               bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                               labels=[1, 2, 3, 4, 5])

housing["income_cat"].value_counts()
```

```
[12]: 3    7236
      2    6581
      4    3639
      5    2362
      1     822
      Name: income_cat, dtype: int64
```

```
[13]: housing["income_cat"].hist()
```

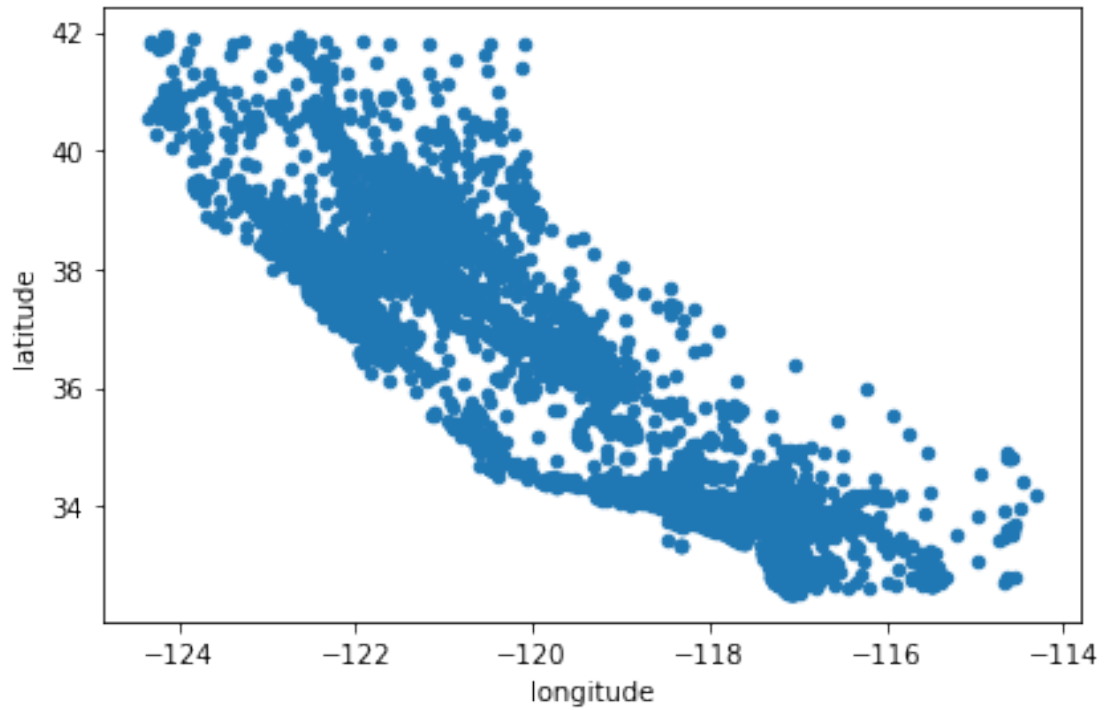
```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x12b769450>
```



Next let's visualize the household incomes based on latitude & longitude coordinates

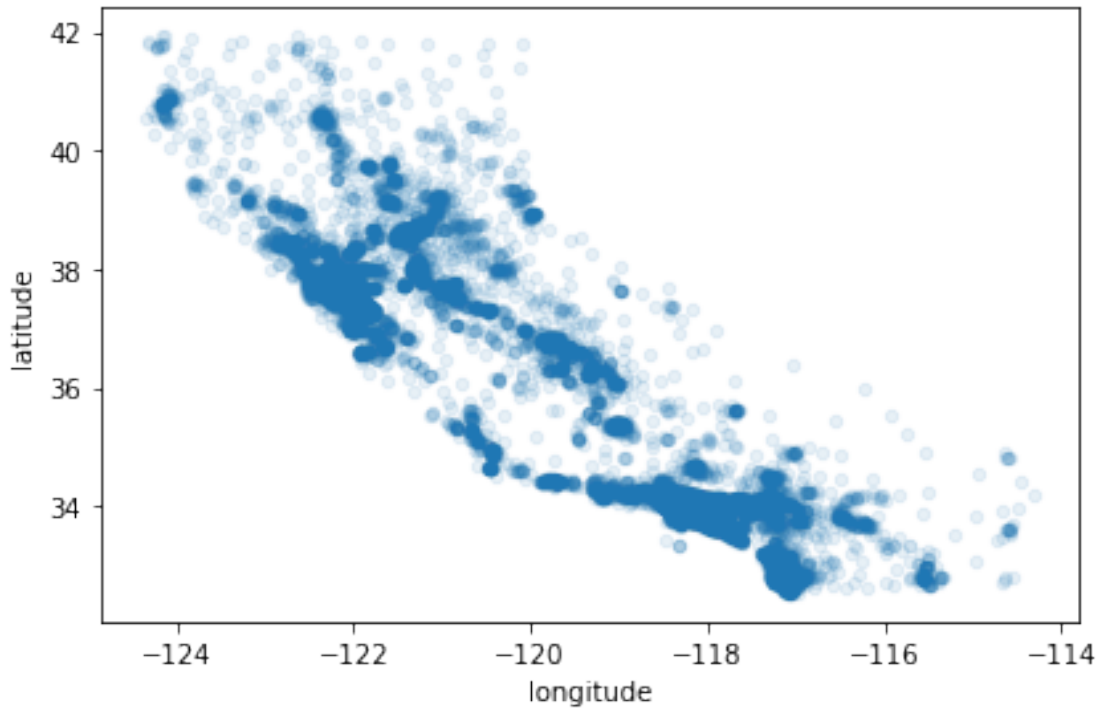
```
[14]: ## here's a not so interesting way plotting it  
housing.plot(kind="scatter", x="longitude", y="latitude")  
save_fig("bad_visualization_plot")
```

Saving figure bad_visualization_plot



```
[15]: # we can make it look a bit nicer by using the alpha parameter,  
# it simply plots less dense areas lighter.  
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)  
save_fig("better_visualization_plot")
```

Saving figure better_visualization_plot



```
[16]: # A more interesting plot is to color code (heatmap) the dots
# based on income. The code below achieves this

# load an image of california
images_path = os.path.join('.', "images")
os.makedirs(images_path, exist_ok=True)
filename = "california.png"

import matplotlib.image as mpimg
california_img=mpimg.imread(os.path.join(images_path, filename))
ax = housing.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
                  s=housing['population']/100, label="Population",
                  c="median_house_value", cmap=plt.get_cmap("jet"),
                  colorbar=False, alpha=0.4,
                  )

# overlay the california map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(california_img, extent=[-124.55, -113.80, 32.45, 42.05], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)
```

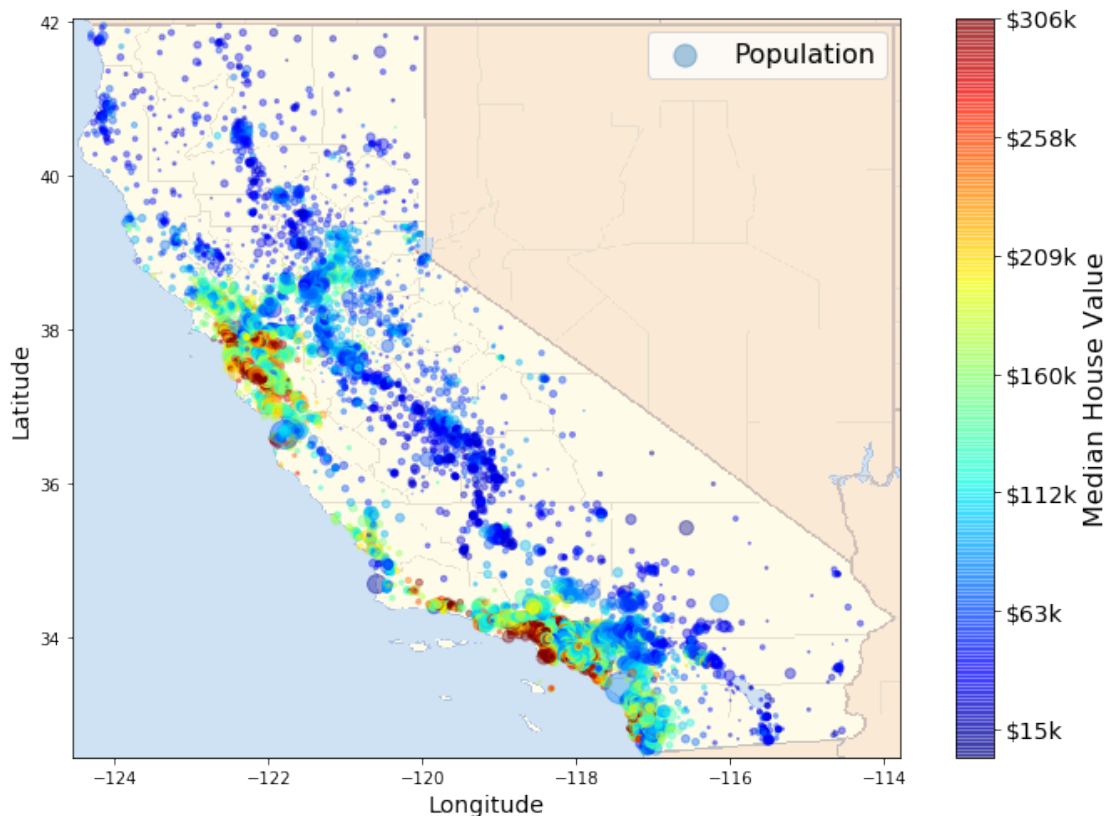
```

# setting up heatmap colors based on median_house_value feature
prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cb = plt.colorbar()
cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values],
    ↪fontsize=14)
cb.set_label('Median House Value', fontsize=16)

plt.legend(fontsize=16)
save_fig("california_housing_prices_plot")
plt.show()

```

Saving figure california_housing_prices_plot



Not suprisingly, the most expensive houses are concentrated around the San Francisco/Los Angeles areas.

Up until now we have only visualized feature histograms and basic statistics.

When developing machine learning models the predictiveness of a feature for a particular target of intrest is what's important.

It may be that only a few features are useful for the target at hand, or features may need to be

augmented by applying certain transformations.

None the less we can explore this using correlation matrices.

```
[17]: corr_matrix = housing.corr()
```

```
[18]: corr_matrix
```

```
[18]:
```

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.000000	-0.924664	-0.108197	0.044568	
latitude	-0.924664	1.000000	0.011173	-0.036100	
housing_median_age	-0.108197	0.011173	1.000000	-0.361262	
total_rooms	0.044568	-0.036100	-0.361262	1.000000	
total_bedrooms	0.069608	-0.066983	-0.320451	0.930380	
population	0.099773	-0.108785	-0.296244	0.857126	
households	0.055310	-0.071035	-0.302916	0.918484	
median_income	-0.015176	-0.079809	-0.119034	0.198050	
median_house_value	-0.045967	-0.144160	0.105623	0.134153	

	total_bedrooms	population	households	median_income	\
longitude	0.069608	0.099773	0.055310	-0.015176	
latitude	-0.066983	-0.108785	-0.071035	-0.079809	
housing_median_age	-0.320451	-0.296244	-0.302916	-0.119034	
total_rooms	0.930380	0.857126	0.918484	0.198050	
total_bedrooms	1.000000	0.877747	0.979728	-0.007723	
population	0.877747	1.000000	0.907222	0.004834	
households	0.979728	0.907222	1.000000	0.013033	
median_income	-0.007723	0.004834	0.013033	1.000000	
median_house_value	0.049686	-0.024650	0.065843	0.688075	

	median_house_value
longitude	-0.045967
latitude	-0.144160
housing_median_age	0.105623
total_rooms	0.134153
total_bedrooms	0.049686
population	-0.024650
households	0.065843
median_income	0.688075
median_house_value	1.000000

```
[19]: # for example if the target is "median_house_value", most correlated features
      ↪ can be sorted
      # which happens to be "median_income". This also intuitively makes sense.
      corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[19]: median_house_value    1.000000
      median_income       0.688075
```

```

total_rooms      0.134153
housing_median_age 0.105623
households       0.065843
total_bedrooms   0.049686
population       -0.024650
longitude        -0.045967
latitude         -0.144160
Name: median_house_value, dtype: float64

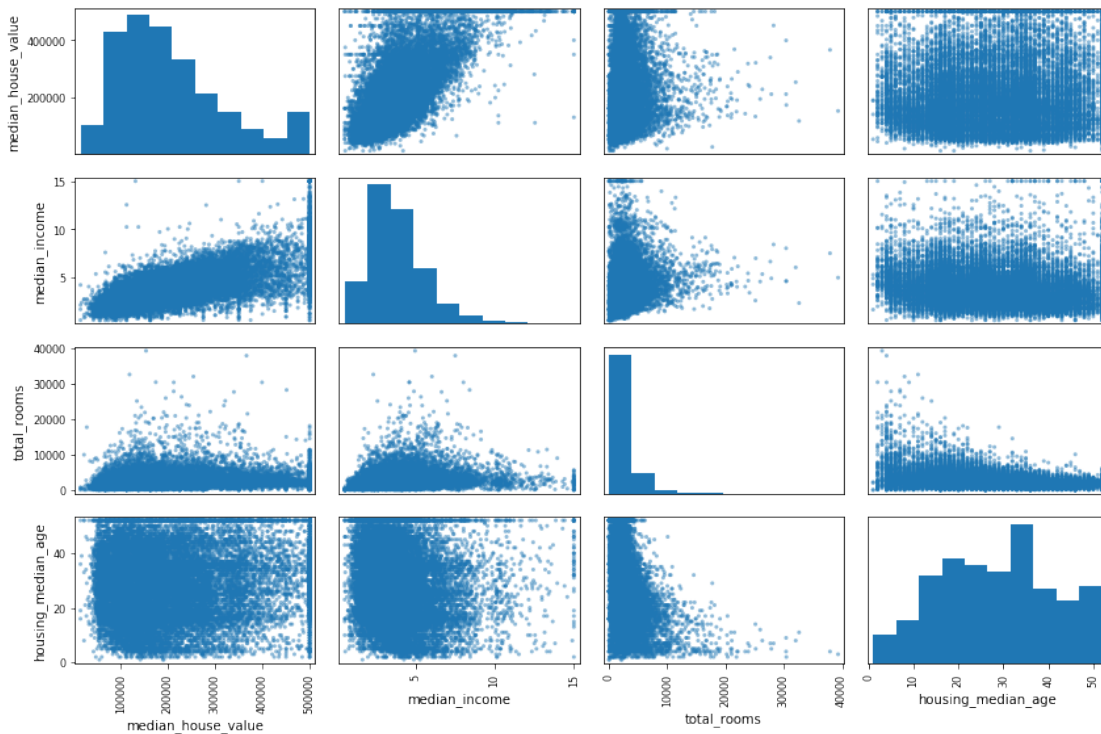
```

```

[20]: # the correlation matrix for different attributes/features can also be plotted
# some features may show a positive correlation/negative correlation or
# it may turn out to be completely random!
from pandas.plotting import scatter_matrix
attributes = ["median_house_value", "median_income", "total_rooms",
             "housing_median_age"]
scatter_matrix(housing[attributes], figsize=(12, 8))
save_fig("scatter_matrix_plot")

```

Saving figure scatter_matrix_plot



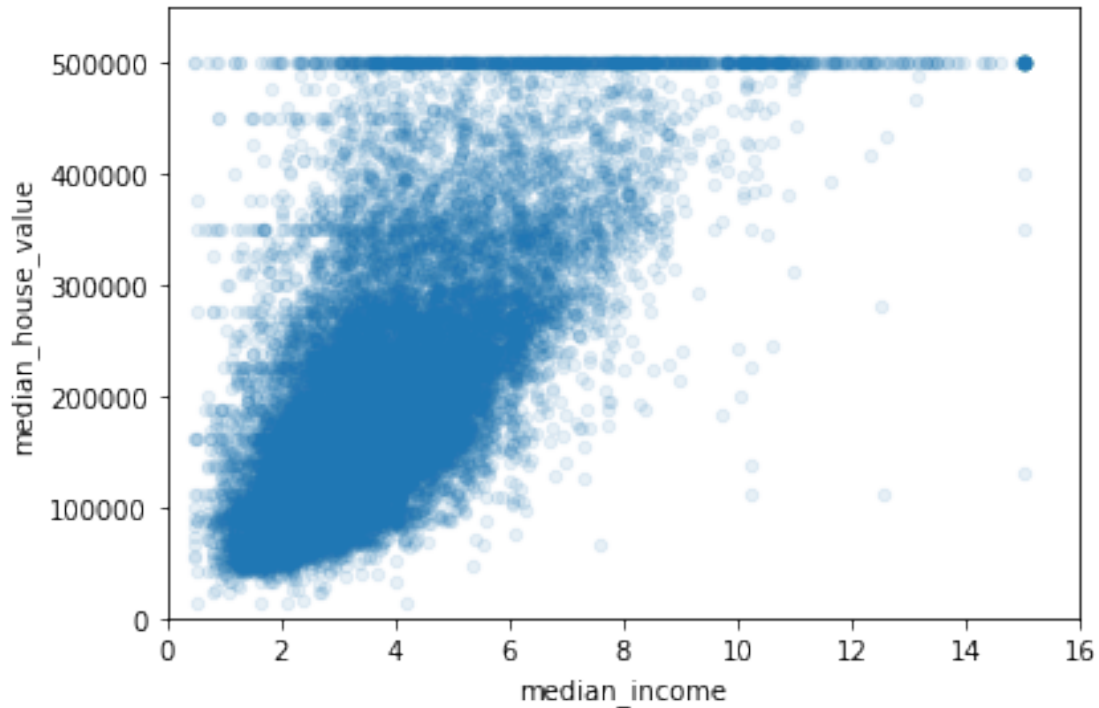
```

[21]: # median income vs median house value plot plot 2 in the first row of top figure
housing.plot(kind="scatter", x="median_income", y="median_house_value",
             alpha=0.1)

```

```
plt.axis([0, 16, 0, 550000])
save_fig("income_vs_house_value_scatterplot")
```

Saving figure income_vs_house_value_scatterplot



0.5.1 Augmenting Features

New features can be created by combining different columns from our data set.

- $\text{rooms_per_household} = \text{total_rooms} / \text{households}$
- $\text{bedrooms_per_room} = \text{total_bedrooms} / \text{total_rooms}$
- etc.

```
[22]: housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"]=housing["population"]/housing["households"]
```

```
[23]: # obtain new correlations
corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[23]: median_house_value    1.000000
      median_income       0.688075
      rooms_per_household  0.151948
```

```

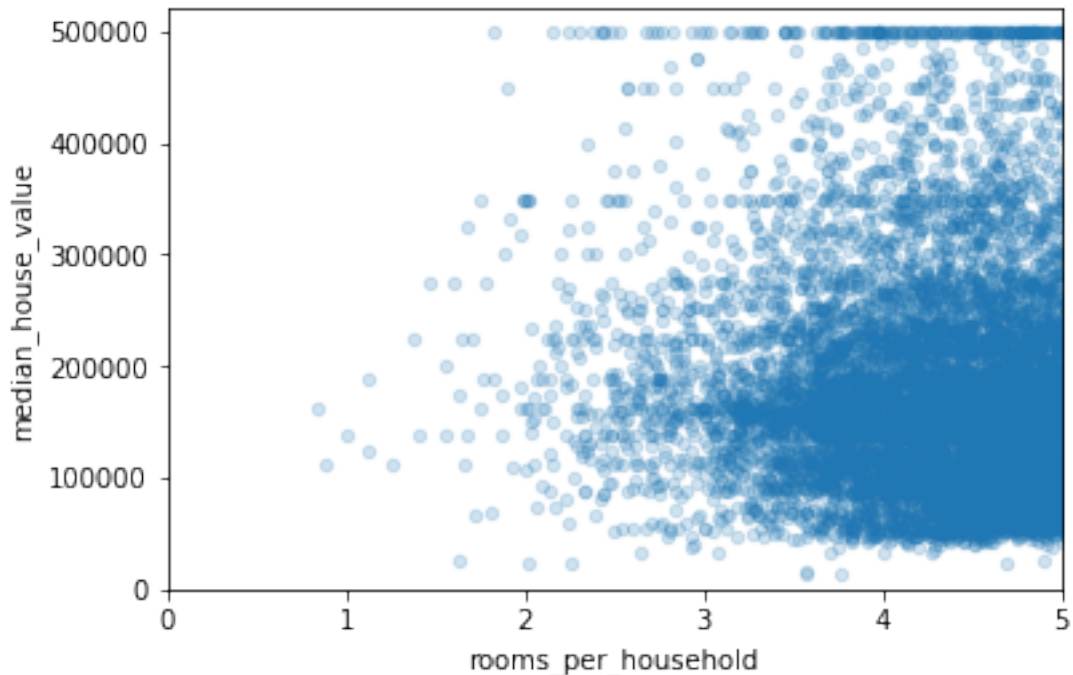
total_rooms          0.134153
housing_median_age    0.105623
households           0.065843
total_bedrooms       0.049686
population_per_household -0.023737
population           -0.024650
longitude            -0.045967
latitude             -0.144160
bedrooms_per_room    -0.255880
Name: median_house_value, dtype: float64

```

```

[24]: housing.plot(kind="scatter", x="rooms_per_household", y="median_house_value",
                alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()

```



```

[25]: housing.describe()

```

```

[25]:
count    longitude    latitude    housing_median_age    total_rooms  \
count    20640.000000    20640.000000    20640.000000    20640.000000
mean      -119.569704     35.631861      28.639486     2635.763081
std         2.003532      2.135952      12.585558     2181.615252
min       -124.350000     32.540000       1.000000       2.000000
25%       -121.800000     33.930000      18.000000     1447.750000

```

50%	-118.490000	34.260000	29.000000	2127.000000
75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value	rooms_per_household	bedrooms_per_room \
count	20640.000000	20640.000000	20433.000000
mean	206855.816909	5.429000	0.213039
std	115395.615874	2.474173	0.057983
min	14999.000000	0.846154	0.100000
25%	119600.000000	4.440716	0.175427
50%	179700.000000	5.229129	0.203162
75%	264725.000000	6.052381	0.239821
max	500001.000000	141.909091	1.000000

	population_per_household
count	20640.000000
mean	3.070655
std	10.386050
min	0.692308
25%	2.429741
50%	2.818116
75%	3.282261
max	1243.333333

0.6 Preparing Dastaset for ML

Once we've visualized the data, and have a certain understanding of how the data looks like. It's time to clean!

Most of your time will be spent on this step, although the datasets used in this project are relatively nice and clean... it could get real dirty.

After having cleaned your dataset you're aiming for: - train set - test set

In some cases you might also have a validation set as well for tuning hyperparameters (don't worry if you're not familiar with this term yet..)

In supervised learning setting your train set and test set should contain (**feature**, **target**) tuples. - **feature**: is the input to your model - **target**: is the ground truth label - when target is categorical

the task is a classification task - when target is floating point the task is a regression task

We will make use of **scikit-learn** python package for preprocessing.

Scikit learn is pretty well documented and if you get confused at any point simply look up the function/object!

```
[26]: from sklearn.model_selection import StratifiedShuffleSplit
# let's first start by creating our train and test sets
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
#print(type(split))
for train_index, test_index in split.split(housing, housing["income_cat"]):
    #print(train_index, test_index)
    train_set = housing.loc[train_index]
    test_set = housing.loc[test_index]
```

```
[27]: train_set.head()
```

```
[27]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
17606	-121.89	37.29	38.0	1568.0	351.0	
18632	-121.93	37.05	14.0	679.0	108.0	
14650	-117.20	32.77	31.0	1952.0	471.0	
3230	-119.61	36.31	25.0	1847.0	371.0	
3555	-118.59	34.23	17.0	6592.0	1525.0	

	population	households	median_income	median_house_value	\
17606	710.0	339.0	2.7042	286600.0	
18632	306.0	113.0	6.4214	340600.0	
14650	936.0	462.0	2.8621	196900.0	
3230	1460.0	353.0	1.8839	46300.0	
3555	4459.0	1463.0	3.0347	254500.0	

	ocean_proximity	income_cat	rooms_per_household	bedrooms_per_room	\
17606	<1H OCEAN	2	4.625369	0.223852	
18632	<1H OCEAN	5	6.008850	0.159057	
14650	NEAR OCEAN	2	4.225108	0.241291	
3230	INLAND	2	5.232295	0.200866	
3555	<1H OCEAN	3	4.505810	0.231341	

	population_per_household
17606	2.094395
18632	2.707965
14650	2.025974
3230	4.135977
3555	3.047847

```
[28]: # Saves the modified training data in the housing variable
```

```

housing = train_set.drop("median_house_value", axis=1) # drop labels for
↳ training set features

# the input to the model
↳ should not contain the true label
# Keep track for later
housing_labels = train_set["median_house_value"].copy()

```

0.6.1 Dealing With Incomplete Data

```

[29]: # have you noticed when looking at the dataframe summary certain rows
# contained null values? we can't just leave them as nulls and expect our
# model to handle them for us...
sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
sample_incomplete_rows

```

```

[29]:
   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
4629    -118.30    34.07                18.0        3759.0           NaN
6068    -117.86    34.01                16.0        4632.0           NaN
17923   -121.97    37.35                30.0        1955.0           NaN
13656   -117.30    34.05                 6.0        2155.0           NaN
19252   -122.79    38.48                 7.0        6837.0           NaN

   population  households  median_income  ocean_proximity  income_cat  \
4629      3296.0      1462.0         2.2708    <1H OCEAN           2
6068      3038.0       727.0         5.1762    <1H OCEAN           4
17923       999.0       386.0         4.6328    <1H OCEAN           4
13656      1039.0       391.0         1.6675      INLAND           2
19252      3468.0      1405.0         3.1662    <1H OCEAN           3

   rooms_per_household  bedrooms_per_room  population_per_household
4629             2.571135              NaN             2.254446
6068             6.371389              NaN             4.178817
17923            5.064767              NaN             2.588083
13656            5.511509              NaN             2.657289
19252            4.866192              NaN             2.468327

```

```

[30]: sample_incomplete_rows.dropna(subset=["total_bedrooms"]) # option 1: simply
↳ drop rows that have null values

```

```

[30]: Empty DataFrame
Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms,
population, households, median_income, ocean_proximity, income_cat,
rooms_per_household, bedrooms_per_room, population_per_household]
Index: []

```

```
[31]: sample_incomplete_rows.drop("total_bedrooms", axis=1)      # option 2: drop
      ↪ the complete feature
```

```
[31]:      longitude  latitude  housing_median_age  total_rooms  population  \
4629      -118.30    34.07             18.0        3759.0        3296.0
6068      -117.86    34.01             16.0        4632.0        3038.0
17923     -121.97    37.35             30.0        1955.0         999.0
13656     -117.30    34.05              6.0        2155.0       1039.0
19252     -122.79    38.48              7.0        6837.0       3468.0

      households  median_income  ocean_proximity  income_cat  \
4629         1462.0         2.2708      <1H OCEAN          2
6068          727.0         5.1762      <1H OCEAN          4
17923         386.0         4.6328      <1H OCEAN          4
13656         391.0         1.6675        INLAND          2
19252        1405.0         3.1662      <1H OCEAN          3

      rooms_per_household  bedrooms_per_room  population_per_household
4629             2.571135                NaN             2.254446
6068             6.371389                NaN             4.178817
17923            5.064767                NaN             2.588083
13656            5.511509                NaN             2.657289
19252            4.866192                NaN             2.468327
```

```
[32]: median = housing["total_bedrooms"].median()
      sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) # option
      ↪ 3: replace na values with median values
      sample_incomplete_rows
```

```
[32]:      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
4629      -118.30    34.07             18.0        3759.0            433.0
6068      -117.86    34.01             16.0        4632.0            433.0
17923     -121.97    37.35             30.0        1955.0            433.0
13656     -117.30    34.05              6.0        2155.0            433.0
19252     -122.79    38.48              7.0        6837.0            433.0

      population  households  median_income  ocean_proximity  income_cat  \
4629         3296.0         1462.0         2.2708      <1H OCEAN          2
6068         3038.0          727.0         5.1762      <1H OCEAN          4
17923          999.0          386.0         4.6328      <1H OCEAN          4
13656         1039.0          391.0         1.6675        INLAND          2
19252         3468.0         1405.0         3.1662      <1H OCEAN          3

      rooms_per_household  bedrooms_per_room  population_per_household
4629             2.571135                NaN             2.254446
6068             6.371389                NaN             4.178817
17923            5.064767                NaN             2.588083
```

13656	5.511509	NaN	2.657289
19252	4.866192	NaN	2.468327

Could you think of another plausible imputation for this dataset? (Not graded)

0.6.2 Prepare Data

```
[33]: # This cell implements the complete pipeline for preparing the data
# using sklearn's TransformerMixins
# Earlier we mentioned different types of features: categorical, and floats.
# In the case of floats we might want to convert them to categories.
# On the other hand categories in which are not already represented as integers
    ↪ must be mapped to integers before
# feeding to the model.

# Additionally, categorical values could either be represented as one-hot
    ↪ vectors or simple as normalized/unnormalized integers.
# Here we encode them using one hot vectors.

from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.base import BaseEstimator, TransformerMixin

imputer = SimpleImputer(strategy="median") # use median imputation for missing
    ↪ values
housing_num = housing.drop("ocean_proximity", axis=1) # remove the categorical
    ↪ feature
# column index
rooms_ix, bedrooms_ix, population_ix, households_ix = 3, 4, 5, 6

#
class AugmentFeatures(BaseEstimator, TransformerMixin):
    """
    implements the previous features we had defined
    housing["rooms_per_household"] = housing["total_rooms"]/
    ↪ housing["households"]
    housing["bedrooms_per_room"] = housing["total_bedrooms"]/
    ↪ housing["total_rooms"]
    housing["population_per_household"] = housing["population"]/
    ↪ housing["households"]
    """
```

```

def __init__(self, add_bedrooms_per_room = True):
    self.add_bedrooms_per_room = add_bedrooms_per_room
def fit(self, X, y=None):
    return self # nothing else to do
def transform(self, X):
    rooms_per_household = X[:, rooms_ix] / X[:, households_ix]
    population_per_household = X[:, population_ix] / X[:, households_ix]
    if self.add_bedrooms_per_room:
        bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
        return np.c_[X, rooms_per_household, population_per_household,
                     bedrooms_per_room]
    else:
        return np.c_[X, rooms_per_household, population_per_household]

attr_adder = AugmentFeatures(add_bedrooms_per_room=False)
housing_extra_attribs = attr_adder.transform(housing.values)

num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', AugmentFeatures()),
    ('std_scaler', StandardScaler()),
])

housing_num_tr = num_pipeline.fit_transform(housing_num)
numerical_features = list(housing_num)
categorical_features = ["ocean_proximity"]

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features),
])

housing_prepared = full_pipeline.fit_transform(housing)

```

```
[34]: housing_prepared[:3]
```

```

[34]: array([[ -1.15604281,  0.77194962,  0.74333089, -0.49323393, -0.44543821,
           -0.63621141, -0.42069842, -0.61493744, -0.95445595, -0.31205452,
            0.19380531, -0.08649871, -0.31205452, -0.08649871,  0.15531753,
            1.          ,  0.          ,  0.          ,  0.          ,  0.          ],
          [ -1.17602483,  0.6596948 , -1.1653172 , -0.90896655, -1.0369278 ,
           -0.99833135, -1.02222705,  1.33645936,  1.89030518,  0.21768338,
           -0.94074539, -0.03353391,  0.21768338, -0.03353391, -0.83628902,
            1.          ,  0.          ,  0.          ,  0.          ,  0.          ],
          [ 1.18684903, -1.34218285,  0.18664186, -0.31365989, -0.15334458,
           -0.43363936, -0.0933178 , -0.5320456 , -0.95445595, -0.46531516,
            0.49916044, -0.09240499, -0.46531516, -0.09240499,  0.4222004 ,

```

```
0.          , 0.          , 0.          , 0.          , 1.          ]])
```

0.6.3 Select a model and train

Once we have prepared the dataset it's time to choose a model.

As our task is to predict the median_house_value (a floating value), regression is well suited for this.

```
[35]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(housing_prepared, housing_labels)

# let's try the full preprocessing pipeline on a few training instances
data = test_set.iloc[:5]
labels = housing_labels.iloc[:5]
data_prepared = full_pipeline.transform(data)

print("Predictions:", lin_reg.predict(data_prepared))
print("Actual labels:", list(labels))
```

```
Predictions: [425717.48517515 267643.98033218 227366.19892733 199614.48287493
161425.25185885]
```

```
Actual labels: [286600.0, 340600.0, 196900.0, 46300.0, 254500.0]
```

```
/usr/local/lib/python3.7/site-
packages/sklearn/compose/_column_transformer.py:430: FutureWarning: Given
feature/column names or counts do not match the ones for the data given during
fit. This will fail from v0.24.
```

```
FutureWarning)
```

We can evaluate our model using certain metrics, a fitting metric for regression is the mean-squared-loss

$$L(\hat{Y}, Y) = \sum_i^N (\hat{y}_i - y_i)^2$$

where \hat{y} is the predicted value, and y is the ground truth label.

```
[36]: from sklearn.metrics import mean_squared_error

preds = lin_reg.predict(housing_prepared)
mse = mean_squared_error(housing_labels, preds)
rmse = np.sqrt(mse)
rmse
```

```
[36]: 67784.32202861732
```

1 TODO: Applying the end-end ML steps to a different dataset.

We will apply what we've learnt to another dataset (airbnb dataset). We will predict airbnb price based on other features.

2 [25 pts] Visualizing Data

2.0.1 [5 pts] Load the data + statistics

- load the dataset
- display the first few rows of the data
- drop the following columns: name, host_id, host_name, last_review
- display a summary of the statistics of the loaded data
- plot histograms for 3 features of your choice

```
[37]: # Loading the dataset
airbnb_data = pd.read_csv('./datasets/airbnb/AB_NYC_2019.csv')

# Displaying the first few rows of the data
airbnb_data.head()
```

```
[37]:      id      name  host_id \
0  2539  Clean & quiet apt home by the park    2787
1  2595      Skylit Midtown Castle    2845
2  3647  THE VILLAGE OF HARLEM...NEW YORK !    4632
3  3831      Cozy Entire Floor of Brownstone    4869
4  5022  Entire Apt: Spacious Studio/Loft by central park    7192

      host_name  neighbourhood_group  neighbourhood  latitude  longitude \
0      John      Brooklyn      Kensington  40.64749  -73.97237
1  Jennifer      Manhattan      Midtown  40.75362  -73.98377
2  Elisabeth      Manhattan      Harlem  40.80902  -73.94190
3  LisaRoxanne      Brooklyn  Clinton Hill  40.68514  -73.95976
4      Laura      Manhattan      East Harlem  40.79851  -73.94399

      room_type  price  minimum_nights  number_of_reviews  last_review \
0  Private room    149              1              9  2018-10-19
1  Entire home/apt    225              1             45  2019-05-21
2  Private room    150              3              0         NaN
3  Entire home/apt     89              1            270  2019-07-05
4  Entire home/apt     80             10              9  2018-11-19

      reviews_per_month  calculated_host_listings_count  availability_365
0              0.21              6              365
1              0.38              2              355
2              NaN              1              365
3              4.64              1              194
4              0.10              1              0
```

```
[38]: airbnb_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   id                                    48895 non-null  int64
 1   name                                48879 non-null  object
 2   host_id                             48895 non-null  int64
 3   host_name                           48874 non-null  object
 4   neighbourhood_group                 48895 non-null  object
 5   neighbourhood                       48895 non-null  object
 6   latitude                           48895 non-null  float64
 7   longitude                           48895 non-null  float64
 8   room_type                           48895 non-null  object
 9   price                               48895 non-null  int64
10  minimum_nights                      48895 non-null  int64
11  number_of_reviews                   48895 non-null  int64
12  last_review                         38843 non-null  object
13  reviews_per_month                  38843 non-null  float64
14  calculated_host_listings_count     48895 non-null  int64
15  availability_365                    48895 non-null  int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB
```

```
[39]: # Draw these columns inplace so it updates the current dataset itself
airbnb_data.drop(['name', 'host_id', 'host_name', 'last_review'], axis=1,
                 ↪inplace=True)
airbnb_data.head()
```

```
[39]:      id  neighbourhood_group  neighbourhood  latitude  longitude  \
0   2539             Brooklyn    Kensington  40.64749   -73.97237
1   2595             Manhattan      Midtown  40.75362   -73.98377
2   3647             Manhattan      Harlem   40.80902   -73.94190
3   3831             Brooklyn  Clinton Hill  40.68514   -73.95976
4   5022             Manhattan    East Harlem  40.79851   -73.94399

      room_type  price  minimum_nights  number_of_reviews  \
0  Private room   149                1                 9
1  Entire home/apt  225                1                45
2  Private room   150                3                 0
3  Entire home/apt   89                1               270
4  Entire home/apt   80               10                 9

      reviews_per_month  calculated_host_listings_count  availability_365
0                0.21                                6                365
```


1	0.38	2	355
2	NaN	1	365
3	4.64	1	194
4	0.10	1	0

```
[40]: # Summary of the statistics of the loaded data:
airbnb_data.describe()
```

```
[40]:
```

	id	latitude	longitude	price	minimum_nights	\
count	4.889500e+04	48895.000000	48895.000000	48895.000000	48895.000000	
mean	1.901714e+07	40.728949	-73.952170	152.720687	7.029962	
std	1.098311e+07	0.054530	0.046157	240.154170	20.510550	
min	2.539000e+03	40.499790	-74.244420	0.000000	1.000000	
25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000	
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000	
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000	
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000	

	number_of_reviews	reviews_per_month	calculated_host_listings_count	\
count	48895.000000	38843.000000	48895.000000	
mean	23.274466	1.373221	7.143982	
std	44.550582	1.680442	32.952519	
min	0.000000	0.010000	1.000000	
25%	1.000000	0.190000	1.000000	
50%	5.000000	0.720000	1.000000	
75%	24.000000	2.020000	2.000000	
max	629.000000	58.500000	327.000000	

	availability_365
count	48895.000000
mean	112.781327
std	131.622289
min	0.000000
25%	0.000000
50%	45.000000
75%	227.000000
max	365.000000

```
[41]: # 3 graphs of features
fig, ax = plt.subplots(3)
fig.subplots_adjust(hspace=0.5)
fig.set_figheight(12)
fig.set_figwidth(8)

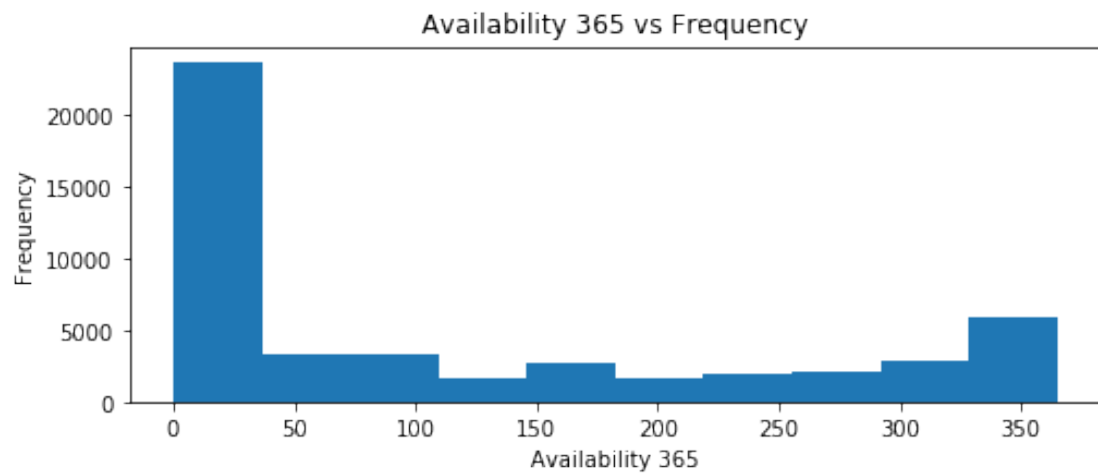
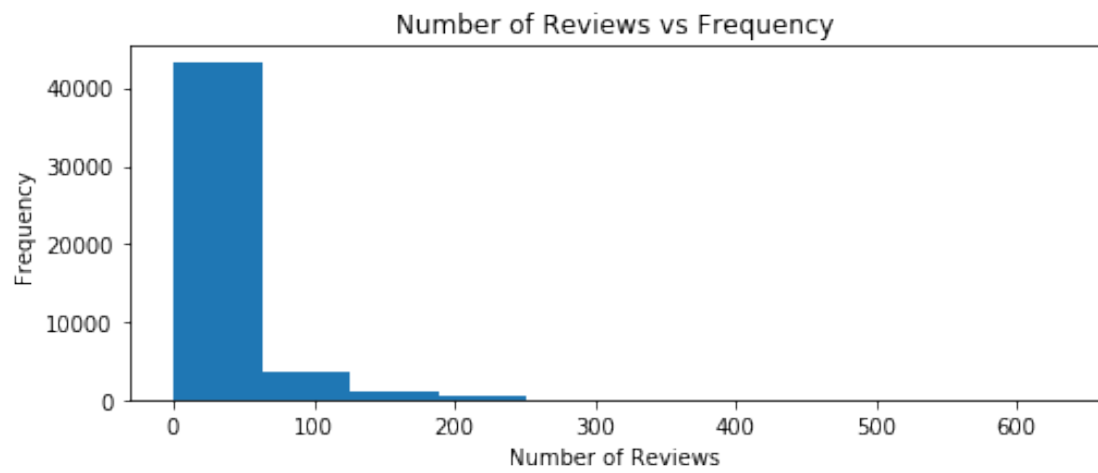
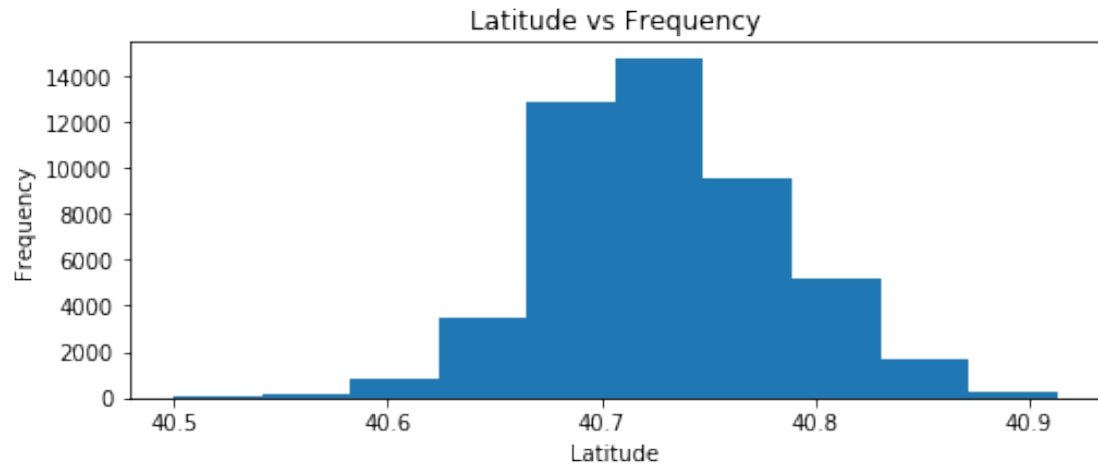
ax[0].hist(airbnb_data["latitude"])
ax[0].set_title('Latitude vs Frequency')
ax[0].set_xlabel('Latitude')
```

```
ax[0].set_ylabel('Frequency')

ax[1].hist(airbnb_data["number_of_reviews"])
ax[1].set_title('Number of Reviews vs Frequency')
ax[1].set_xlabel('Number of Reviews')
ax[1].set_ylabel('Frequency')

ax[2].hist(airbnb_data["availability_365"])
ax[2].set_title('Availability 365 vs Frequency')
ax[2].set_xlabel('Availability 365')
ax[2].set_ylabel('Frequency')

plt.show()
```



2.0.2 [5 pts] Plot total number_of_reviews per neighbourhood_group

```
[42]: data_subset = airbnb_data[['number_of_reviews', 'neighbourhood_group']]
data_subset.head()

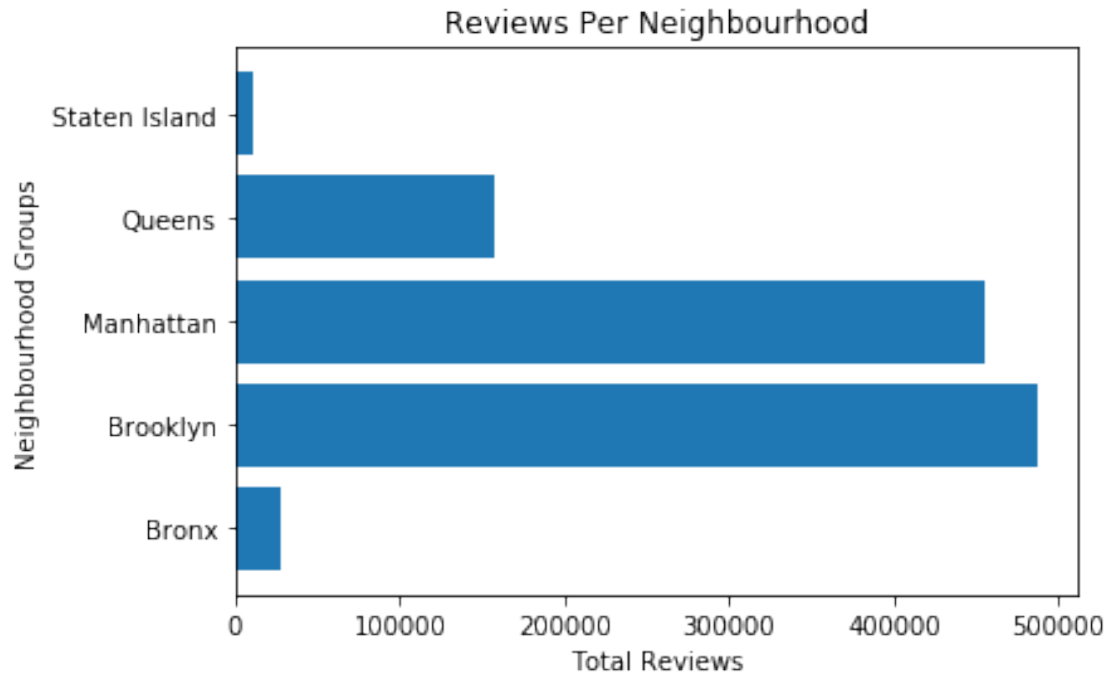
# Use sum to calculate `total` number of reviews. Like a SQL group by.
reviews_per_neighbourhood_group = data_subset.groupby(['neighbourhood_group']).
    ↪sum().reset_index()
reviews_per_neighbourhood_group.columns =
    ↪['neighbourhood_group', 'total_reviews']

print(reviews_per_neighbourhood_group)

neighbourhood_group = reviews_per_neighbourhood_group['neighbourhood_group'].
    ↪tolist()
total_reviews = reviews_per_neighbourhood_group['total_reviews'].tolist()

fig, ax = plt.subplots()
ax.barh(neighbourhood_group, total_reviews)
ax.set_title('Reviews Per Neighbourhood')
ax.set_xlabel('Total Reviews')
ax.set_ylabel('Neighbourhood Groups')
plt.show()
```

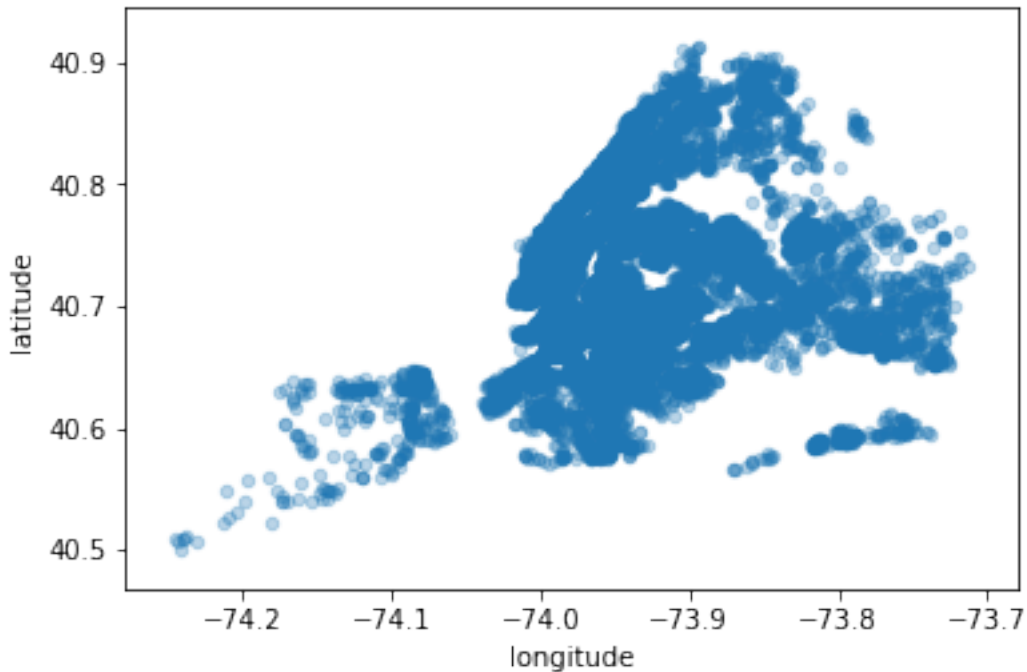
	neighbourhood_group	total_reviews
0	Bronx	28371
1	Brooklyn	486574
2	Manhattan	454569
3	Queens	156950
4	Staten Island	11541



2.0.3 [5 pts] Plot map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

```
[43]: airbnb_data.plot(kind='scatter',x='longitude',y='latitude',alpha=0.3)
```

```
[43]: <matplotlib.axes._subplots.AxesSubplot at 0x13152b050>
```



```
[44]: # Finding coordinates for boundary
latitude_min = airbnb_data['latitude'].min()
latitude_max = airbnb_data['latitude'].max()
longitude_min = airbnb_data['longitude'].min()
longitude_max = airbnb_data['longitude'].max()
BBox = (longitude_min, longitude_max, latitude_min, latitude_max)
print(BBox)
```

```
(-74.24441999999999, -73.71299, 40.499790000000004, 40.913059999999994)
```

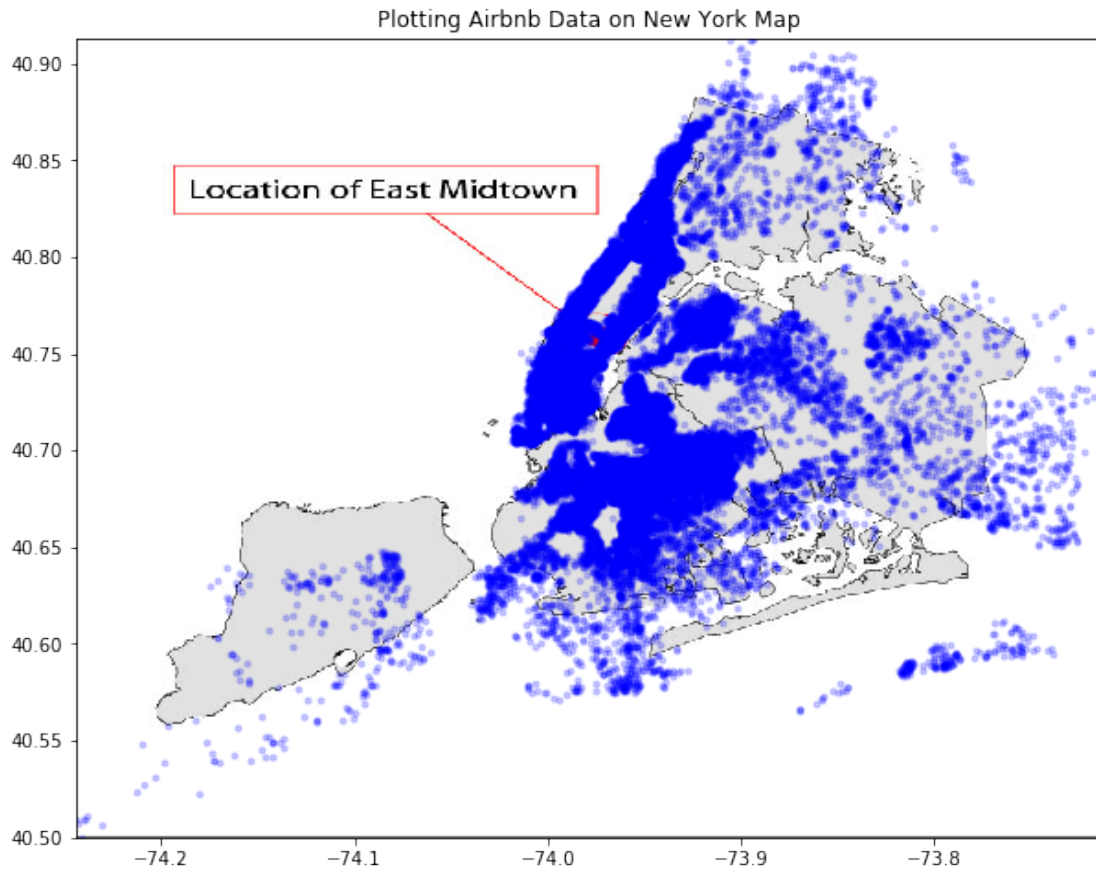
```
[45]: # load an image of greater new york
images_path = os.path.join('./', "images")
os.makedirs(images_path, exist_ok=True)
filename = "greater_new_york.jpg"

import matplotlib.image as mpimg
newyork_img=mpimg.imread(os.path.join(images_path, filename))

fig, ax = plt.subplots(figsize = (12,8))
ax.scatter(airbnb_data.longitude, airbnb_data.latitude, zorder=1, alpha=0.2,
    ↪c='b', s=10)
ax.set_title('Plotting Airbnb Data on New York Map')
ax.set_xlim(BBox[0],BBox[1])
ax.set_ylim(BBox[2],BBox[3])
```

```
ax.imshow(newyork_img, zorder=0, extent = BBox)
```

[45]: <matplotlib.image.AxesImage at 0x1308950d0>



2.0.4 [5 pts] Plot average price of room types who have availability greater than 180 days.

```
[46]: availability = airbnb_data[airbnb_data['availability_365'] > 180]
availability = availability[['room_type', 'price']]
availability = availability.groupby(['room_type']).mean().reset_index()

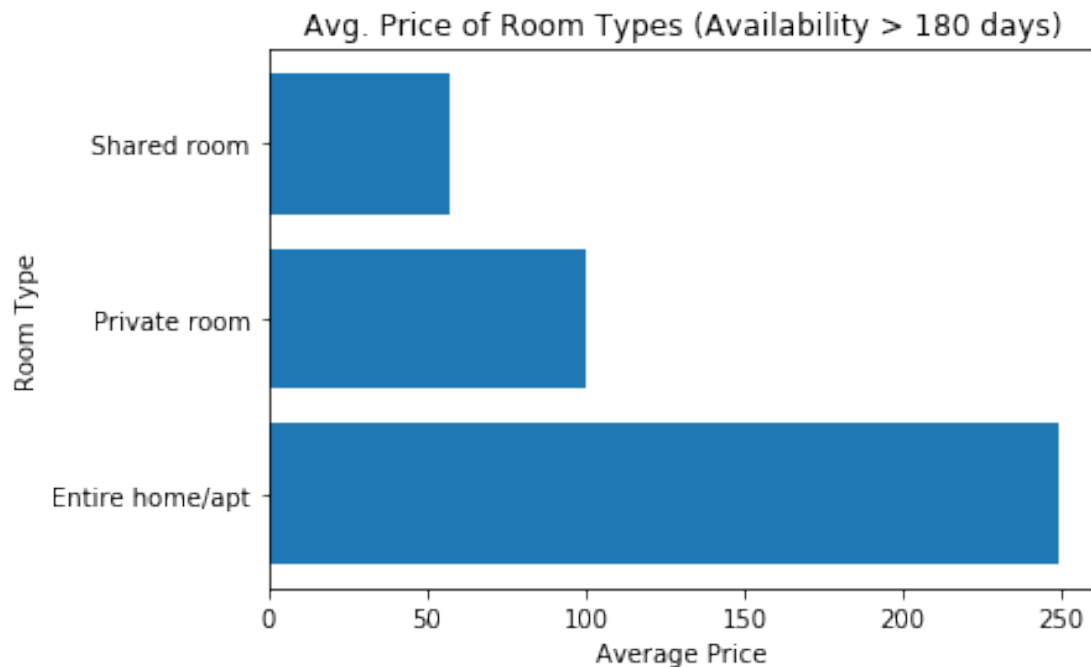
print(availability)

room_types = availability['room_type'].tolist()
prices = availability['price'].tolist()

fig, ax = plt.subplots()
ax.barh(room_types, prices)
ax.set_title('Avg. Price of Room Types (Availability > 180 days)')
```

```
ax.set_xlabel('Average Price')
ax.set_ylabel('Room Type')
plt.show()
```

	room_type	price
0	Entire home/apt	248.870817
1	Private room	100.028192
2	Shared room	56.941909



2.0.5 [5 pts] Plot correlation matrix

- which features have positive correlation?
- which features have negative correlation?

```
[47]: airbnb_data.columns
```

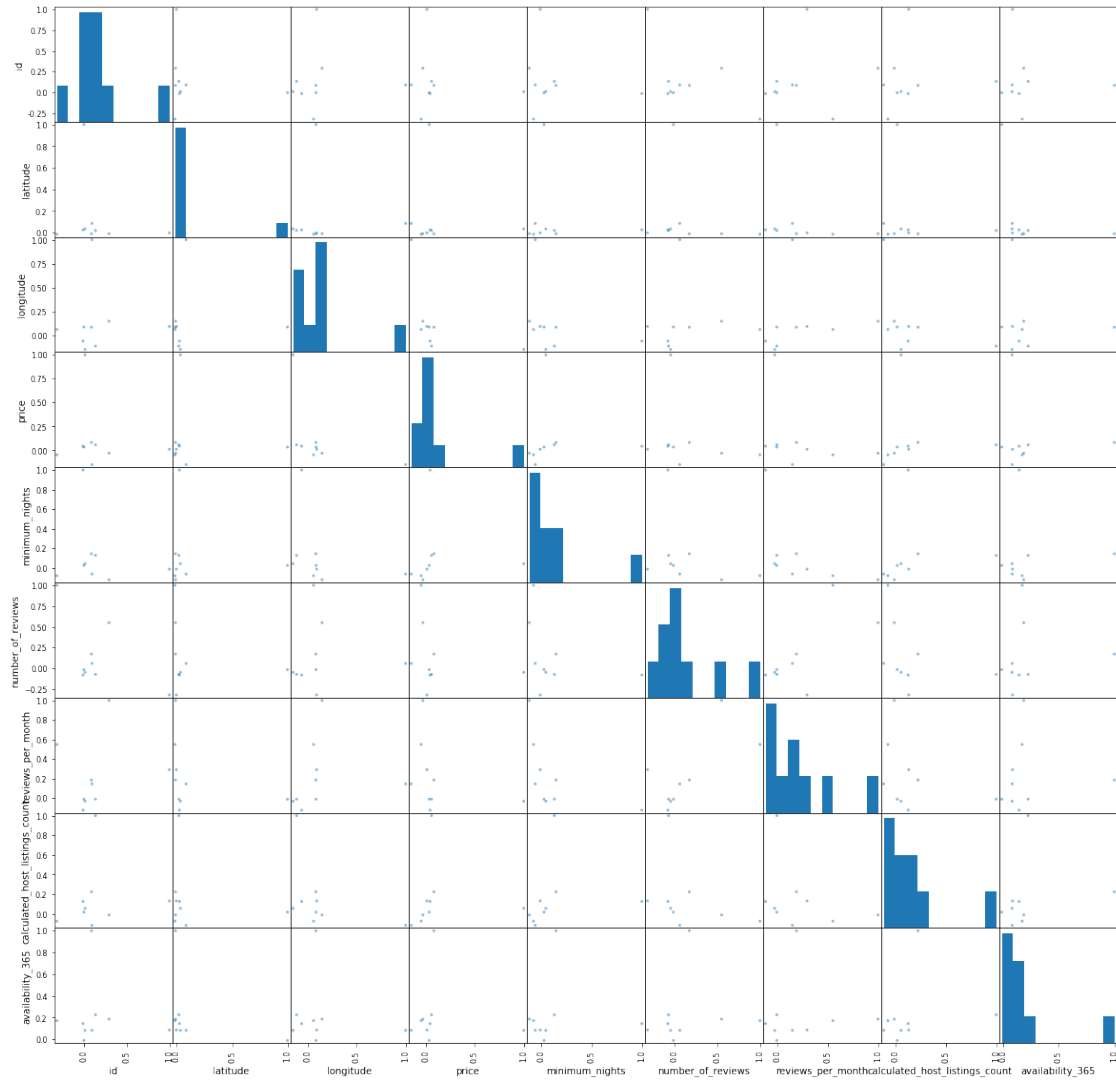
```
[47]: Index(['id', 'neighbourhood_group', 'neighbourhood', 'latitude', 'longitude',
          'room_type', 'price', 'minimum_nights', 'number_of_reviews',
          'reviews_per_month', 'calculated_host_listings_count',
          'availability_365'],
          dtype='object')
```

```
[48]: # No categorical, remove latitude and longitude
features=['price', 'minimum_nights', 'number_of_reviews', 'reviews_per_month',
          'calculated_host_listings_count', 'availability_365']
```



```
scatter_matrix(airbnb_data[features], figsize=(12,12))
```

```
[48]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x130e8f910>,
<matplotlib.axes._subplots.AxesSubplot object at 0x1308b7950>,
<matplotlib.axes._subplots.AxesSubplot object at 0x130a32790>,
<matplotlib.axes._subplots.AxesSubplot object at 0x13092f890>,
<matplotlib.axes._subplots.AxesSubplot object at 0x1309cf3d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12b16fa90>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x12b435f10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12b5af8d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12b3c2550>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12bc40210>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12bc03d50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12afea290>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x12b875190>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12c7a04d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12b5c54d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12b76a610>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12ed31b10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x130940e90>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x12c8bab50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12fd3ced0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x13091fb90>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12c7fcf10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12ec5bbd0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12e98cf50>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x12ece5c10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x13002af90>,
<matplotlib.axes._subplots.AxesSubplot object at 0x130057c50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12b911fd0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12be34c90>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12bfc14d0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x12c515cd0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12c556510>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12c58bd10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12c5cc550>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12c602d50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x12c642590>]],
dtype=object)
```



```
[49]: corr = airbnb_data.corr()
      corr["price"].sort_values(ascending=False)
```

```
[49]: price          1.000000
      availability_365  0.081829
      calculated_host_listings_count  0.057472
      minimum_nights    0.042799
      latitude          0.033939
      id               0.010619
      reviews_per_month -0.030608
      number_of_reviews -0.047954
      longitude        -0.150019
      Name: price, dtype: float64
```

Based on the values seen above from the correlation of all input features with the target variable (price), the relationship can be split into two - namely positive correlation and negative correlation.

Positively Correlated Features:

- availability_365 0.081829
- calculated_host_listings_count 0.057472
- minimum_nights 0.042799
- latitude 0.033939
- id 0.010619

The most positively correlated feature (with price) in the input data is **availability_365**.

Negatively Correlated Features:

- reviews_per_month -0.030608
- number_of_reviews -0.047954
- longitude -0.150019

The most negatively correlated feature (with price) in the input data is **longitude**.

3 [25 pts] Prepare the Data

3.0.1 [5 pts] Augment the dataframe with two other features which you think would be useful

```
[50]: airbnb_data['price_by_reviews'] = airbnb_data['price']/  
      ↪airbnb_data['number_of_reviews']  
airbnb_data['reviews_per_host_listings'] = airbnb_data['number_of_reviews']/  
      ↪airbnb_data['calculated_host_listings_count']
```

```
[51]: airbnb_data.head(5)
```

```
[51]:
```

	id	neighbourhood_group	neighbourhood	latitude	longitude	\
0	2539	Brooklyn	Kensington	40.64749	-73.97237	
1	2595	Manhattan	Midtown	40.75362	-73.98377	
2	3647	Manhattan	Harlem	40.80902	-73.94190	
3	3831	Brooklyn	Clinton Hill	40.68514	-73.95976	
4	5022	Manhattan	East Harlem	40.79851	-73.94399	

	room_type	price	minimum_nights	number_of_reviews	\
0	Private room	149	1	9	
1	Entire home/apt	225	1	45	
2	Private room	150	3	0	
3	Entire home/apt	89	1	270	
4	Entire home/apt	80	10	9	

	reviews_per_month	calculated_host_listings_count	availability_365	\
0	0.21	6	365	

1	0.38	2	355
2	NaN	1	365
3	4.64	1	194
4	0.10	1	0

	price_by_reviews	reviews_per_host_listings
0	16.555556	1.5
1	5.000000	22.5
2	inf	0.0
3	0.329630	270.0
4	8.888889	9.0

3.0.2 [5 pts] Impute any missing feature with a method of your choice, and briefly discuss why you chose this imputation method

```
[52]: airbnb_data[airbnb_data.isnull().any(axis=1)].head(3)
```

```
[52]:      id neighbourhood_group neighbourhood latitude longitude \
2    3647             Manhattan      Harlem  40.80902  -73.94190
19   7750             Manhattan  East Harlem  40.79685  -73.94872
26   8700             Manhattan      Inwood  40.86754  -73.92639

      room_type price minimum_nights number_of_reviews \
2    Private room    150             3                0
19  Entire home/apt    190             7                0
26    Private room     80             4                0

      reviews_per_month calculated_host_listings_count availability_365 \
2                  NaN                      1                365
19                 NaN                      2                249
26                 NaN                      1                 0

      price_by_reviews reviews_per_host_listings
2                  inf                0.0
19                 inf                0.0
26                 inf                0.0
```

```
[53]: airbnb_data.isna().head(3)
```

```
[53]:      id neighbourhood_group neighbourhood latitude longitude room_type \
0  False                False            False      False      False  False
1  False                False            False      False      False  False
2  False                False            False      False      False  False

      price minimum_nights number_of_reviews reviews_per_month \
0  False                False            False                False
1  False                False            False                False
```

	False	False	False	True
	calculated_host_listings_count	availability_365	price_by_reviews	\
0	False	False	False	
1	False	False	False	
2	False	False	False	

	False
	reviews_per_host_listings
0	False
1	False
2	False

```
[54]: # Price by reviews is not a very logical feature, and additionally, theres too
      ↪many
      # records with number_of_reviews = 0, and just dropping those rows would
      ↪greatly
      # reduce the size of this dataset. Therefore, I decided to drop this entire
      ↪feature.
      airbnb_data = airbnb_data.drop('price_by_reviews', axis=1)

      # For reviews_per_month, I decided to find the median value and use it to
      ↪substitute
      # all the NaN values in the 'reviews_per_month' column. Instead of losing rows
      ↪of
      # data, I felt like this was more statistically significant and made more
      ↪sense, and a
      # somewhat decent assumption to make, although it will skew the loss value
      ↪significantly.
      reviews_per_month_median = airbnb_data['reviews_per_month'].median()
      airbnb_data['reviews_per_month'].fillna(reviews_per_month_median, inplace=True)

      # After the above imputations, if there were still any rows with NaNs (about
      ↪5-10 rows),
      # I decided to drop them altogether since we had dealt with most of the NaNs in
      ↪a logical
      # way prior to this.
      airbnb_data.dropna(inplace=True)
```

3.0.3 [10 pts] Code complete data pipeline using sklearn mixins

```
[55]: # Drop ID since it is a unique feature (and quite irrelevant)
      # One hot-encoding vector will be massive = equal to number of unique IDs,
      ↪leading to a highly sparse matrix
      airbnb_data = airbnb_data.drop('id', axis=1)
```

```
[56]: print(len(list(airbnb_data['neighbourhood_group'].unique())))
      # Primary source of feature expansion in model when this categorical feature is
      # → one-hot encoded
      print(len(list(airbnb_data['neighbourhood'].unique())))
      print(len(list(airbnb_data['room_type'].unique())))
```

```
5
221
3
```

```
[57]: # Create a copy of the output features to predict
      labels = airbnb_data['price'].copy()
      airbnb_data = airbnb_data.drop('price', axis=1)
```

```
[58]: from sklearn.impute import SimpleImputer
      from sklearn.compose import ColumnTransformer

      from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler
      from sklearn.preprocessing import OneHotEncoder

      from sklearn.base import BaseEstimator, TransformerMixin

      categorical_features = ['neighbourhood_group', 'neighbourhood', 'room_type']
      airbnb_no_categorical = airbnb_data.drop(categorical_features, axis=1)

      num_pipeline = Pipeline([
          ('imputer', SimpleImputer(strategy="median")),
          ('std_scaler', StandardScaler()),
      ])

      numerical_features = list(airbnb_no_categorical)

      full_pipeline = ColumnTransformer([
          ("num", num_pipeline, numerical_features),
          ("cat", OneHotEncoder(), categorical_features),
      ])

      airbnb_prepared = full_pipeline.fit_transform(airbnb_data)
```

```
[59]: airbnb_prepared[:3]
```

```
[59]: <3x237 sparse matrix of type '<class 'numpy.float64'>'
      with 33 stored elements in Compressed Sparse Row format>
```

Primary source for the expansion/creation of so many features comes from one-hot encoding, but particularly, from the **neighbourhood** feature, which has **220 unique values**, resulting in a really

sparse one hot encoded matrix.

3.0.4 [5 pts] Set aside 20% of the data as test set (80% train, 20% test).

```
[60]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(airbnb_prepared, labels,
↳ test_size=0.2, random_state=23)
```

```
[61]: X_train.shape
```

```
[61]: (39116, 237)
```

```
[62]: X_test.shape
```

```
[62]: (9779, 237)
```

```
[63]: y_train.shape
```

```
[63]: (39116,)
```

```
[64]: y_test.shape
```

```
[64]: (9779,)
```

4 [15 pts] Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using MSE. Provide both test and train set MSE values.

```
[65]: lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
```

```
[65]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
[66]: # Test set MSE
preds = lin_reg.predict(X_test)
mse = mean_squared_error(y_test, preds)
print("Test MSE:",mse)
```

Test MSE: 46378.827120361464

```
[67]: # Train set MSE
preds = lin_reg.predict(X_train)
mse = mean_squared_error(y_train, preds)
print("Train MSE:",mse)
```

Train MSE: 52102.518239250894